

Pokémon and Machine Learning: Mega-evolve Prediction

December 11, 2025
Machine Learning Operations Final Project

*Group 7 – Bradley Stoller, Samuel Martinez Koss,
Xigang Zhang, and Zhiwei Guo*



Agenda

1 Dataset Introduction and EDA Insights

Brief overview of the data and key insights from the EDA

2 Pipeline Staging and Evaluation Framework

Review of pipeline and metrics used for evaluation

3 Modeling Methodology

Summary of AutoML and Mlflow results

4 Model Monitoring Approach

Overview of model monitoring setup and usage method

5 Data Change

Discussion of data changes for drift detection confirmation

6 Results and Demo

Review of all results and model monitoring drift detection demo

Agenda

1 Dataset Introduction and EDA Insights

Brief overview of the data and key insights from the EDA

2 Pipeline Staging and Evaluation Framework

Review of pipeline and metrics used for evaluation

3 Modeling Methodology

Summary of AutoML and Mlflow results

4 Model Monitoring Approach

Overview of model monitoring setup and usage method

5 Data Change

Discussion of data changes for drift detection confirmation

6 Results and Demo

Review of all results and model monitoring drift detection demo

Introduction – Why Pokémon?

Data Background

- From Niantic, the company developed Pokémon Go
- It contains 21 features (variables) spanning 721 Pokémon (N=721)

Analysis Objective

- Strongest primary type (species)
- Features that correlated with Mega-evolve
- Build a **predictive model** to determine whether a Pokémon can **Mega-evolve**

Analysis Significance

- For the Pokémon game itself, provide data-driven insights about Pokémon stats to guide future balance and design
- Enable Pokémon-related companies to use data-driven decision making for designing new Pokémon and potentially increasing revenue
- For Pokémon players (like me!), this helps us understand which Pokémon are most likely to succeed in battle



What is a Mega Evolution?



- Increased attack/defense
- Bonus skills
- Exclusive features

High-level Dataset Introduction: Pokémon Characters

721

Total
Rows

21

Features
Available

13

Numerical
Features

7

Categorical
Features

2

Prediction
Classes

Dataset Overview:

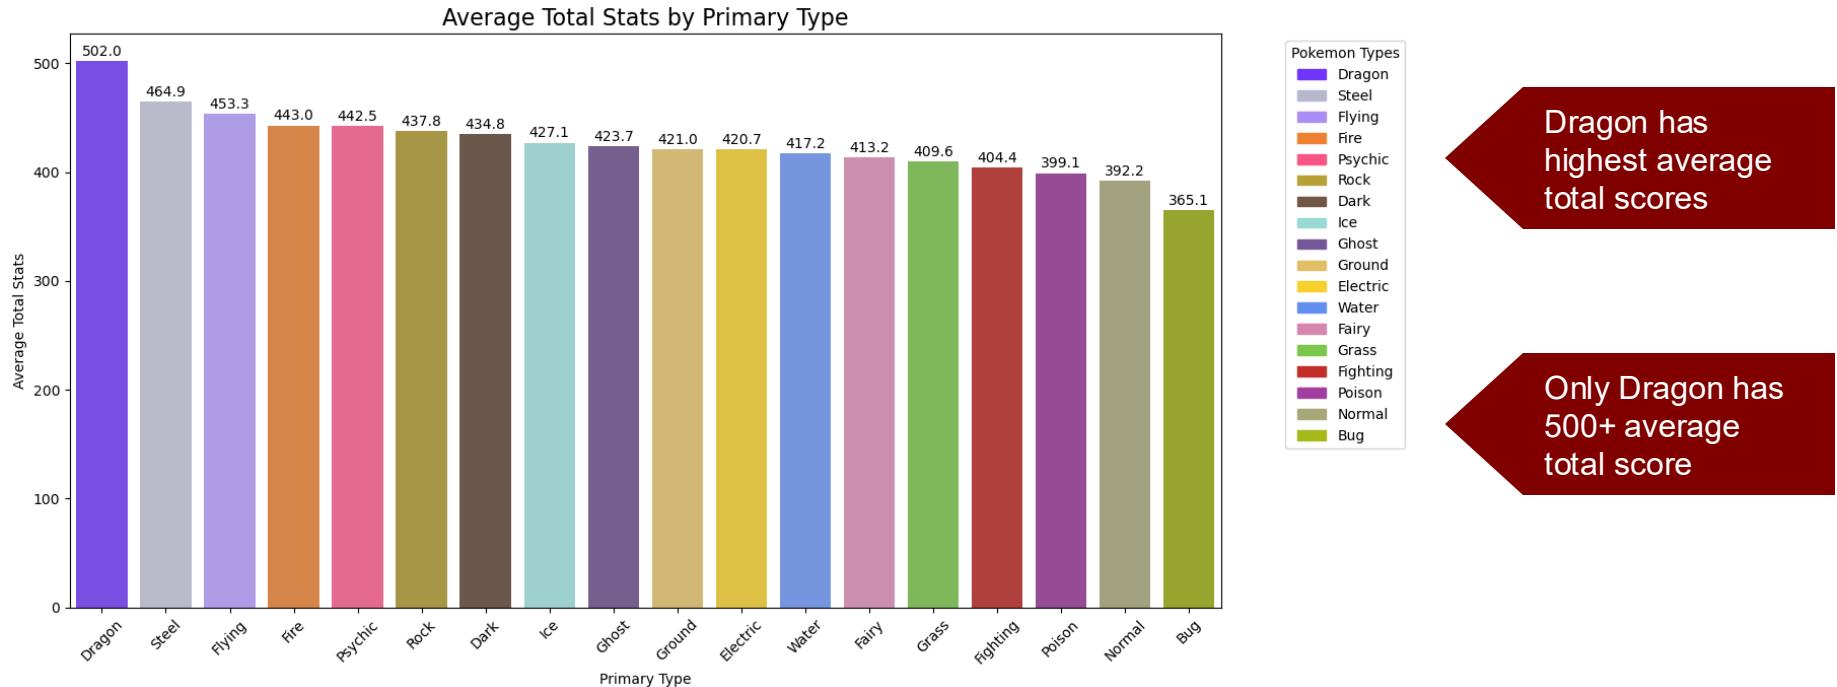
- This dataset focuses on the stats and features of Pokémon from Generations 1 through 6
- It contains 21 features (variables) spanning 721 Pokémon

Target Variable:

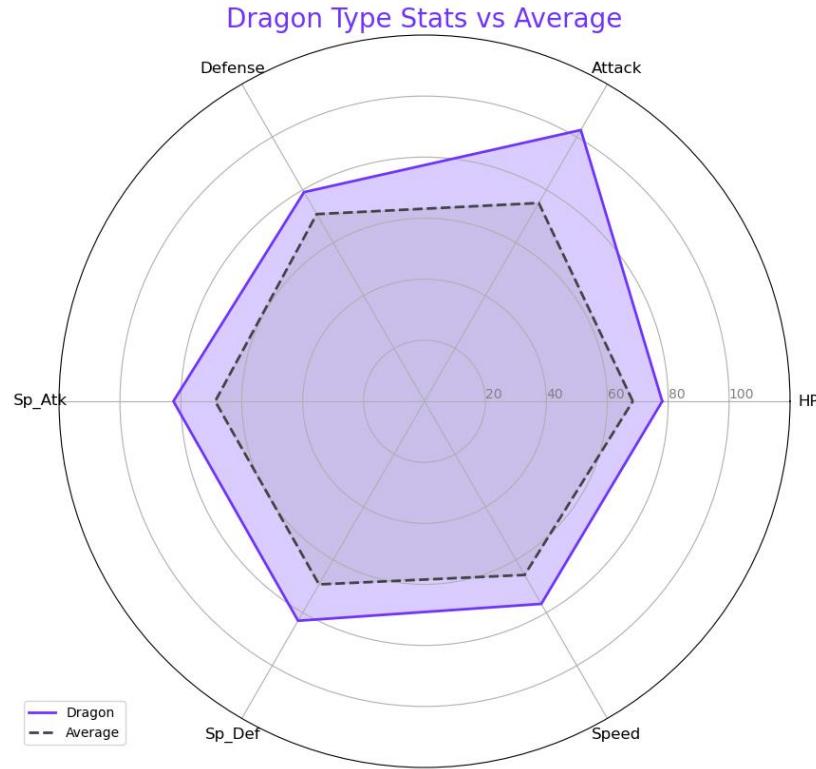
hasMegaEvolution



Overall, Dragon is the most powerful Primary Type across the dataset



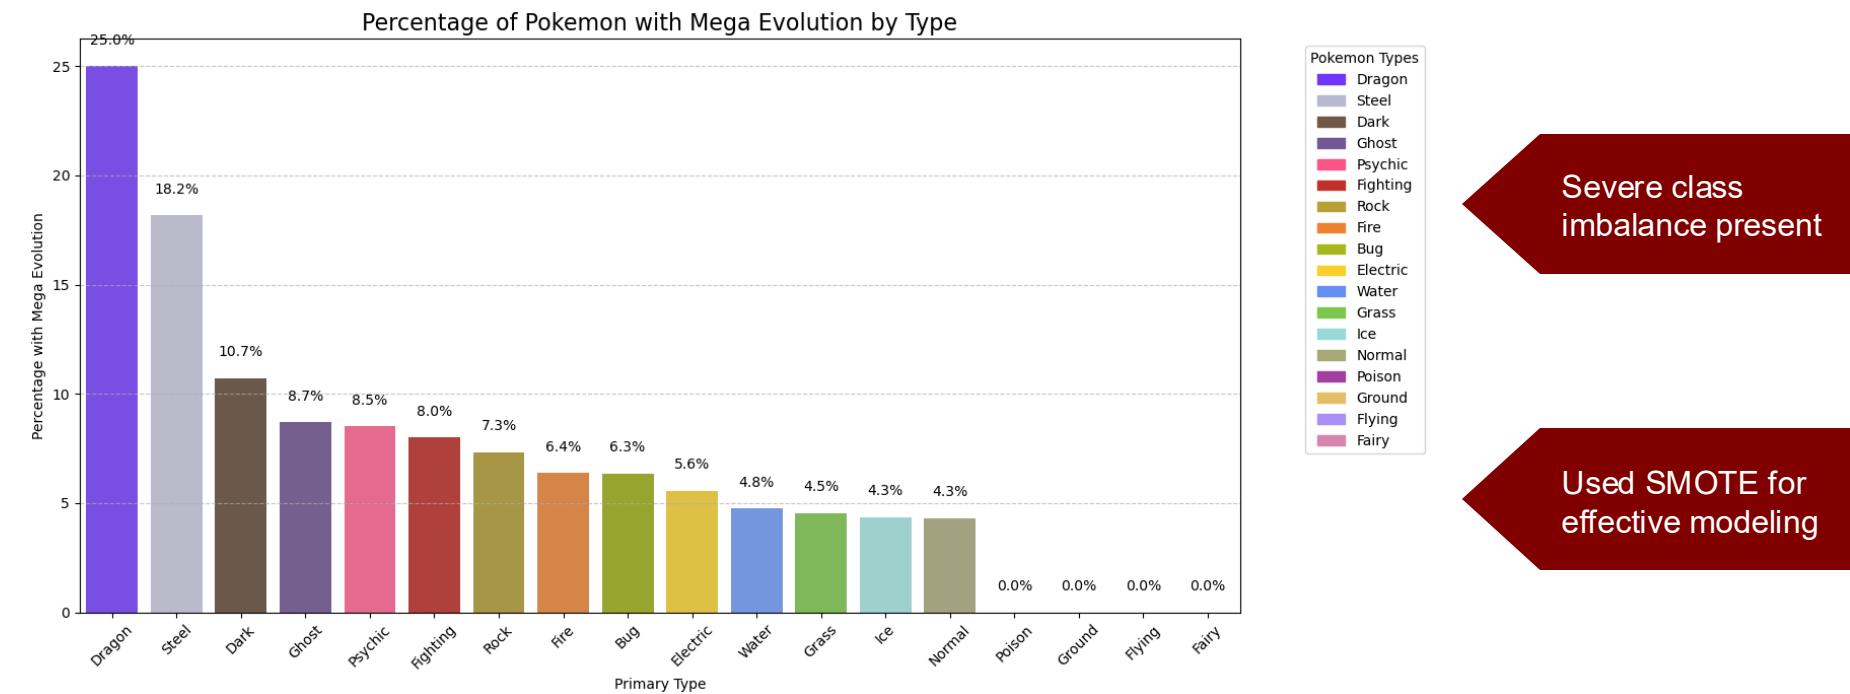
Specifically, Dragon's stats is higher than all other types' stat



Dragon's average
stats is higher than
all other characters

Attack is the most
significant one

However, beyond Dragon, we see that the data presents a class imbalance



Agenda

1

Dataset Introduction and EDA Insights

Brief overview of the data and key insights from the EDA

2

Pipeline Staging and Evaluation Framework

Review of pipeline and metrics used for evaluation

3

Modeling Methodology

Summary of AutoML and Mlflow results

4

Model Monitoring Approach

Overview of model monitoring setup and usage method

5

Data Change

Discussion of data changes for drift detection confirmation

6

Results and Demo

Review of all results and model monitoring drift detection demo

For the evaluation metric, we chose accuracy since it is the most wholistic

		Predicted
		Evolution No Evolution
Actual	Evolution	✓ TP
	No Evolution	✗ FN
Actual	No Evolution	✗ FP
	No Evolution	✓ TN

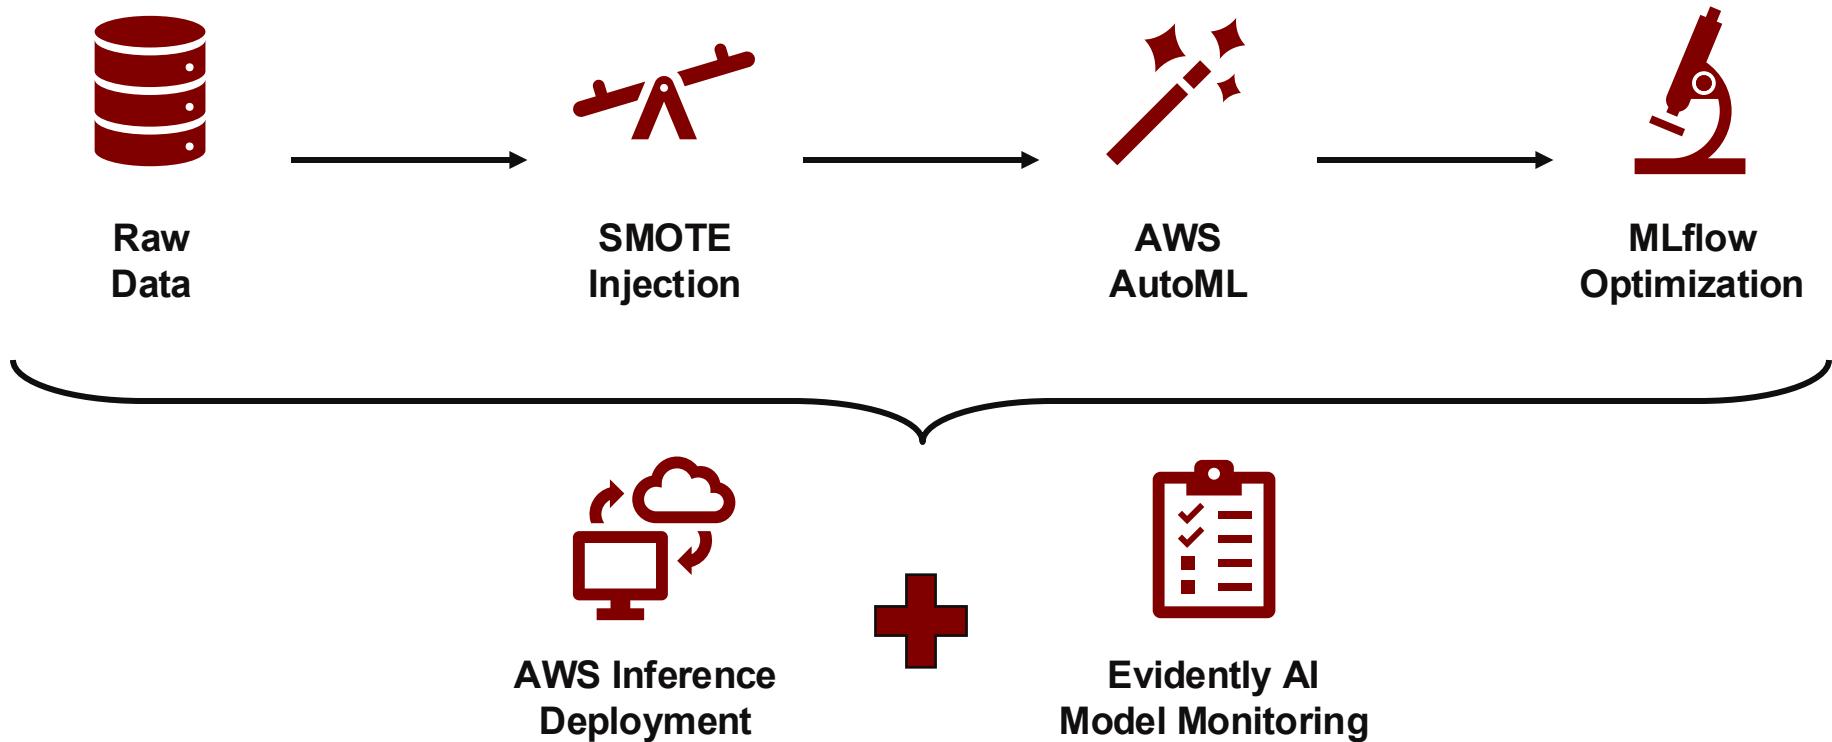
Accuracy:

$$\frac{(TP + TN)}{(TP + TN + FP + FN)}$$



- Balanced representation of both classes after SMOTE
- Easy to interpret with regards to character evolution
- Appropriate given the equal cost of FPs and FNs

For prediction, we used a multi-stage pipeline with AutoML, MLflow, and AWS



Specifically, our MLflow pipeline tests 10 models via random search

n_estimators

2 – 7

learning_rate

0.01 – 0.1

max_depth

2 – 4

min_samples_split

5 – 20

R
A
N
D
O
M

S
E
A
R
C
H

- 1 n_estimators: 2, learning_rate: 0.054, max_depth: 3, min_samples_split: 11
- 2 n_estimators: 7, learning_rate: 0.080, max_depth: 2, min_samples_split: 6
- 3 n_estimators: 5, learning_rate: 0.086, max_depth: 3, min_samples_split: 15
- 4 n_estimators: 7, learning_rate: 0.022, max_depth: 3, min_samples_split: 11
- 5 n_estimators: 5, learning_rate: 0.010, max_depth: 2, min_samples_split: 16
- 6 n_estimators: 5, learning_rate: 0.092, max_depth: 3, min_samples_split: 11
- 7 n_estimators: 3, learning_rate: 0.065, max_depth: 2, min_samples_split: 18
- 8 n_estimators: 2, learning_rate: 0.093, max_depth: 2, min_samples_split: 14
- 9 n_estimators: 2, learning_rate: 0.045, max_depth: 3, min_samples_split: 6
- 10 n_estimators: 7, learning_rate: 0.099, max_depth: 3, min_samples_split: 16

Agenda

1 Dataset Introduction and EDA Insights

Brief overview of the data and key insights from the EDA

2 Pipeline Staging and Evaluation Framework

Review of pipeline and metrics used for evaluation

3 Modeling Methodology

Summary of AutoML and Mlflow results

4 Model Monitoring Approach

Overview of model monitoring setup and usage method

5 Data Change

Discussion of data changes for drift detection confirmation

6 Results and Demo

Review of all results and model monitoring drift detection demo

From AWS AutoML, we found that gradient boosting was the best algorithm



Best Candidate Model

- **Accuracy:** 97%
- **Objective Metric:** Validation Accuracy
- **Detected Algorithm Family:** XG-Boost/Gradient Boosting

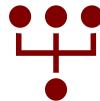
Best Candidate Model Hyperparameters

- **max_depth:** 4
- **eta:** 0.66
- **num_round:** 364
- **colsample_bytree:** 0.96
- **gamma:** 0.00035
- **lambda:** 0.96

After MLflow experimentation, our deployed model achieves 88% train accuracy

88%

Train Accuracy



n_estimators:
7



learning_rate:
0.099



max_depth:
3



min_samples_split:
16

Agenda

1 Dataset Introduction and EDA Insights

Brief overview of the data and key insights from the EDA

2 Pipeline Staging and Evaluation Framework

Review of pipeline and metrics used for evaluation

3 Modeling Methodology

Summary of AutoML and Mlflow results

4 Model Monitoring Approach

Overview of model monitoring setup and usage method

5 Data Change

Discussion of data changes for drift detection confirmation

6 Results and Demo

Review of all results and model monitoring drift detection demo

To monitor data and prediction drift, we used Evidently AI via the AWS Suite

SageMaker Model

Evidently AI Script

S3 Record Archive

Monitoring Dashboard



Evidently integrates seamlessly with AWS workflow



S3 centralizes records for reliable audits and traceability

Agenda

1 Dataset Introduction and EDA Insights

Brief overview of the data and key insights from the EDA

2 Pipeline Staging and Evaluation Framework

Review of pipeline and metrics used for evaluation

3 Modeling Methodology

Summary of AutoML and Mlflow results

4 Model Monitoring Approach

Overview of model monitoring setup and usage method

5 Data Change

Discussion of data changes for drift detection confirmation

6 Results and Demo

Review of all results and model monitoring drift detection demo

We made four key changes to X_test to test the model monitoring effectiveness

Change #1:

Swapped column 0 and column 1

Change #2:

Swapped column 2 and column 7

Change #3:

Randomized column 2

Change #4:

Randomized column 4

New Accuracy 74%
(16% decrease)

Agenda

1 Dataset Introduction and EDA Insights

Brief overview of the data and key insights from the EDA

2 Pipeline Staging and Evaluation Framework

Review of pipeline and metrics used for evaluation

3 Modeling Methodology

Summary of AutoML and Mlflow results

4 Model Monitoring Approach

Overview of model monitoring setup and usage method

5 Data Change

Discussion of data changes for drift detection confirmation

6 Results and Demo

Review of all results and model monitoring drift detection demo

Demo – Model Monitoring In Action!

The screenshot shows a Jupyter Notebook interface with the following details:

- Title Bar:** Notebooks > ml_ops_final_project
- Section Header:** ml_ops_final_project
- Description:** Add description
- Last Saved:** Last saved: 1 minute ago
- Code Cell:** 5b. Deploy the native XG-Boost model to AWS for inference
- Code Content:**

```
▶ ▾ ✓ 1 minute ago | Elapsed time: 3s
1 # Set S3 configurations for model saving
2 MODEL_KEY = 'model/model.tar.gz'
3 MODEL_S3_PATH = f's3://{BUCKET_NAME}/{MODEL_KEY}'
4 ENDPOINT_NAME = 'pokemon-model'
5
6 # Save the model
7 xgb_model.save_model('xgboost-model')
8
9 # Convert to tar.gz format (SageMaker requirement)
10 with tarfile.open('model.tar.gz', 'w:gz') as tar:
11     tar.add('xgboost-model')
12     print("Model packaged as model.tar.gz")
13
14 # Upload to S3
15 S3_CLIENT.upload_file(
16     Filename='model.tar.gz',
17     Bucket=BUCKET_NAME,
18     Key=MODEL_KEY
19 )
20 print(f"Model uploaded to: {S3_PATH}")
21
22 xgb_model_sm = XGBoostModel(
23     model_data= MODEL_S3_PATH,
24     role= ROLE_ARN,
25     framework_version= '1.7-1',
26     sagemaker_session= SESSION
27 )
28
29 print("\nDeploying model to SageMaker endpoint...")
30
```
- Kernel:** Python
- Resource Usage:** Python 3.11, 2 vCPU + 4 GiB, 4.00%, 53.99%, 2.00%
- Status:** Ready

Drift Report Comparison: X_test

1	1	0	0	0
Tests	Success	Warning	Fail	Error

All tests ▾

Number of Drifted Features

The drift is detected for 2 out of 21 features. The test threshold is lt=3.

Feature name	Stattest	Drift score	Threshold	Data Drift
Type_2	K-S p_value	0.017	0.05	Detected
Has_Type_2	Z-test p_value	0.002	0.05	Detected
prediction	Z-test p_value	0.163	0.05	Not detected
Type_1	K-S p_value	0.562	0.05	Not detected
Total	K-S p_value	0.053	0.05	Not detected
Speed	K-S p_value	0.217	0.05	Not detected
Sp_Def	K-S p_value	0.08	0.05	Not detected
Sp_Atk	K-S p_value	0.468	0.05	Not detected
Height_m	K-S p_value	0.17	0.05	Not detected
Weight_kg	K-S p_value	0.545	0.05	Not detected
Has_Egg_Group_2	Z-test p_value	0.077	0.05	Not detected
Generation	K-S p_value	0.945	0.05	Not detected
Egg_Group_2	K-S p_value	0.545	0.05	Not detected
Egg_Group_1	K-S p_value	0.999	0.05	Not detected
Defense	K-S p_value	0.056	0.05	Not detected

[Edit report](#)

Drift Report Comparison: X_test_modified

1	0	0	1	0
Tests	Success	Warning	Fail	Error
All tests ▾				
Number of Drifted Features				
①	The drift is detected for 7 out of 21 features. The test threshold is lt=3.			Details
Feature name	Stattest	Drift score	Threshold	Data Drift
prediction	Z-test p_value	0.0	0.05	Detected
Attack	K-S p_value	0.0	0.05	Detected
Type_2	K-S p_value	0.0	0.05	Detected
Type_1	K-S p_value	0.0	0.05	Detected
Total	K-S p_value	0.0	0.05	Detected
Sp_Def	K-S p_value	0.005	0.05	Detected
Has_Type_2	Z-test p_value	0.002	0.05	Detected
Speed	K-S p_value	0.217	0.05	Not detected
Sp_Atk	K-S p_value	0.468	0.05	Not detected
Height_m	K-S p_value	0.17	0.05	Not detected
Has_Egg_Group_2	Z-test p_value	0.077	0.05	Not detected
HP	K-S p_value	0.295	0.05	Not detected
Generation	K-S p_value	0.945	0.05	Not detected
Egg_Group_2	K-S p_value	0.545	0.05	Not detected
Egg_Group_1	K-S p_value	0.000	0.05	Not detected

Questions?