

# Pokémon and Machine Learning: Mega-evolve Prediction

December 11, 2025  
Machine Learning Operations Final Project

*Group 7 – Bradley Stoller, Samuel Martinez Koss,  
Xigang Zhang, and Zhiwei Guo*



# Agenda

## 1 Dataset Introduction and EDA Insights

Brief overview of the data and key insights from the EDA

## 2 Pipeline Staging and Evaluation Framework

Review of pipeline and metrics used for evaluation

## 3 Modeling Methodology

Summary of AutoML and Mlflow results

## 4 Model Monitoring Approach

Overview of model monitoring setup and usage method

## 5 Data Change

Discussion of data changes for drift detection confirmation

## 6 Results and Demo

Review of all results and model monitoring drift detection demo

# Agenda

## 1 Dataset Introduction and EDA Insights

Brief overview of the data and key insights from the EDA

## 2 Pipeline Staging and Evaluation Framework

Review of pipeline and metrics used for evaluation

## 3 Modeling Methodology

Summary of AutoML and Mlflow results

## 4 Model Monitoring Approach

Overview of model monitoring setup and usage method

## 5 Data Change

Discussion of data changes for drift detection confirmation

## 6 Results and Demo

Review of all results and model monitoring drift detection demo

# Introduction – Why Pokémon?

## Data Background

- From Niantic, the company developed Pokémon Go
- It contains 21 features (variables) spanning 721 Pokémon (N=721)

## Analysis Objective

- Strongest primary type (species)
- Features that correlated with Mega-evolve
- Build a **predictive model** to determine whether a Pokémon can **Mega-evolve**

## Analysis Significance

- For the Pokémon game itself, provide data-driven insights about Pokémon stats to guide future balance and design
- Enable Pokémon-related companies to use data-driven decision making for designing new Pokémon and potentially increasing revenue
- For Pokémon players (like me!), this helps us understand which Pokémon are most likely to succeed in battle



# What is a Mega Evolution?



- Increased attack/defense
- Bonus skills
- Exclusive features

# High-level Dataset Introduction: Pokémon Characters

**721**

Total Rows

**21**

Features Available

**13**

Numerical Features

**7**

Categorical Features

**2**

Prediction Classes

## Dataset Overview:


- This dataset focuses on the stats and features of Pokémon from Generations 1 through 6
- It contains 21 features (variables) spanning 721 Pokémon

## Target Variable:

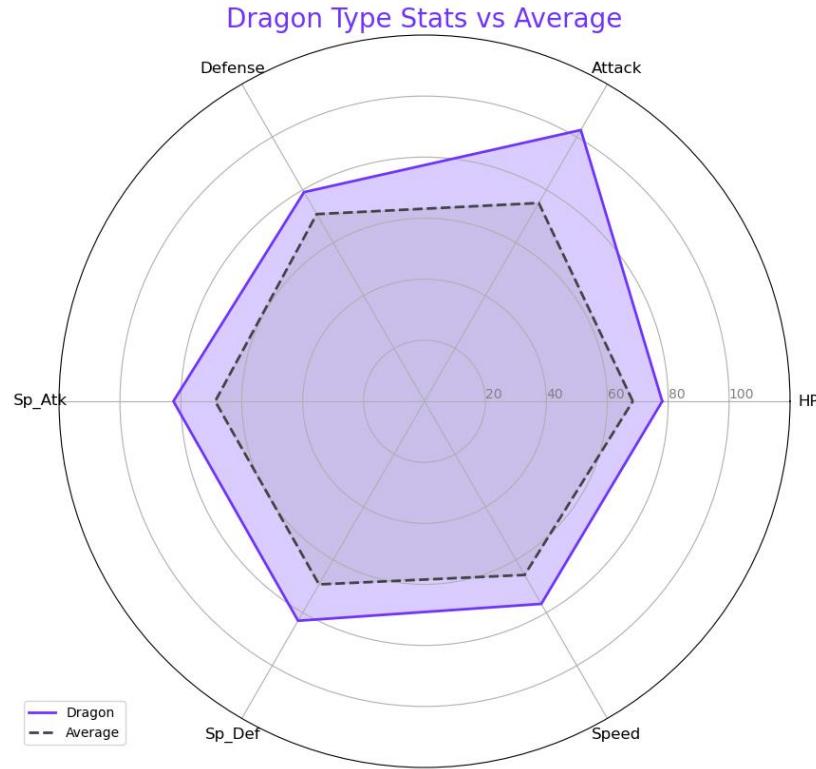
**hasMegaEvolution**



# Overall, Dragon is the most powerful Primary Type across the dataset




# Specifically, Dragon's stats is higher than all other types' stat



Dragon's average  
stats is higher than  
all other characters

Attack is the most  
significant one

# However, beyond Dragon, we see that the data presents a class imbalance



# Agenda

1

## Dataset Introduction and EDA Insights

Brief overview of the data and key insights from the EDA

2

## Pipeline Staging and Evaluation Framework

Review of pipeline and metrics used for evaluation

3

## Modeling Methodology

Summary of AutoML and Mlflow results

4

## Model Monitoring Approach

Overview of model monitoring setup and usage method

5

## Data Change

Discussion of data changes for drift detection confirmation

6

## Results and Demo

Review of all results and model monitoring drift detection demo

# For the evaluation metric, we chose accuracy since it is the most wholistic

		Predicted
		Evolution      No Evolution
Actual	Evolution	✓ TP
	No Evolution	✗ FN
Actual	No Evolution	✗ FP
	Evolution	✓ TN


**Accuracy:**

$$\frac{(TP + TN)}{(TP + TN + FP + FN)}$$



- Balanced representation of both classes after SMOTE
- Easy to interpret with regards to character evolution
- Appropriate given the equal cost of FPs and FNs

For prediction, we used a multi-stage pipeline with AutoML, MLflow, and AWS



# Specifically, our MLflow pipeline tests 10 models via random search

n\_estimators

2 – 7

learning\_rate

0.01 – 0.1

max\_depth

2 – 4

min\_samples\_split

5 – 20

R  
A  
N  
D  
O  
M  
  
S  
E  
A  
R  
C  
H

- 1 n\_estimators: 2, learning\_rate: 0.054, max\_depth: 3, min\_samples\_split: 11
- 2 n\_estimators: 7, learning\_rate: 0.080, max\_depth: 2, min\_samples\_split: 6
- 3 n\_estimators: 5, learning\_rate: 0.086, max\_depth: 3, min\_samples\_split: 15
- 4 n\_estimators: 7, learning\_rate: 0.022, max\_depth: 3, min\_samples\_split: 11
- 5 n\_estimators: 5, learning\_rate: 0.010, max\_depth: 2, min\_samples\_split: 16
- 6 n\_estimators: 5, learning\_rate: 0.092, max\_depth: 3, min\_samples\_split: 11
- 7 n\_estimators: 3, learning\_rate: 0.065, max\_depth: 2, min\_samples\_split: 18
- 8 n\_estimators: 2, learning\_rate: 0.093, max\_depth: 2, min\_samples\_split: 14
- 9 n\_estimators: 2, learning\_rate: 0.045, max\_depth: 3, min\_samples\_split: 6
- 10 n\_estimators: 7, learning\_rate: 0.099, max\_depth: 3, min\_samples\_split: 16

# Agenda

## 1 Dataset Introduction and EDA Insights

Brief overview of the data and key insights from the EDA

## 2 Pipeline Staging and Evaluation Framework

Review of pipeline and metrics used for evaluation

## 3 Modeling Methodology

Summary of AutoML and Mlflow results

## 4 Model Monitoring Approach

Overview of model monitoring setup and usage method

## 5 Data Change

Discussion of data changes for drift detection confirmation

## 6 Results and Demo

Review of all results and model monitoring drift detection demo

From AWS AutoML, we found that gradient boosting was the best algorithm



### Best Candidate Model

- **Accuracy:** 97%
- **Objective Metric:** Validation Accuracy
- **Detected Algorithm Family:** XG-Boost/Gradient Boosting

### Best Candidate Model Hyperparameters

- **max\_depth:** 4
- **eta:** 0.66
- **num\_round:** 364
- **colsample\_bytree:** 0.96
- **gamma:** 0.00035
- **lambda:** 0.96

After MLflow experimentation, our deployed model achieves 88% train accuracy

88%

Train Accuracy



n\_estimators:  
7



learning\_rate:  
0.099



max\_depth:  
3



min\_samples\_split:  
16

# Agenda

## 1 Dataset Introduction and EDA Insights

Brief overview of the data and key insights from the EDA

## 2 Pipeline Staging and Evaluation Framework

Review of pipeline and metrics used for evaluation

## 3 Modeling Methodology

Summary of AutoML and Mlflow results

## 4 Model Monitoring Approach

Overview of model monitoring setup and usage method

## 5 Data Change

Discussion of data changes for drift detection confirmation

## 6 Results and Demo

Review of all results and model monitoring drift detection demo

## To monitor data and prediction drift, we used Evidently AI via the AWS Suite

SageMaker  
Model

Evidently AI  
Script

S3 Record  
Archive

Monitoring  
Dashboard



**Evidently integrates seamlessly with AWS workflow**



**S3 centralizes records for reliable audits and traceability**

# Agenda

## 1 Dataset Introduction and EDA Insights

Brief overview of the data and key insights from the EDA

## 2 Pipeline Staging and Evaluation Framework

Review of pipeline and metrics used for evaluation

## 3 Modeling Methodology

Summary of AutoML and Mlflow results

## 4 Model Monitoring Approach

Overview of model monitoring setup and usage method

## 5 Data Change

Discussion of data changes for drift detection confirmation

## 6 Results and Demo

Review of all results and model monitoring drift detection demo

# We made four key changes to X\_test to test the model monitoring effectiveness

**Change #1:**

Swapped column 0 and column 1

**Change #2:**

Swapped column 2 and column 7

**Change #3:**

Randomized column 2

**Change #4:**

Randomized column 4

**New Accuracy 74%**  
*(16% decrease)*

# Agenda

## 1 Dataset Introduction and EDA Insights

Brief overview of the data and key insights from the EDA

## 2 Pipeline Staging and Evaluation Framework

Review of pipeline and metrics used for evaluation

## 3 Modeling Methodology

Summary of AutoML and Mlflow results

## 4 Model Monitoring Approach

Overview of model monitoring setup and usage method

## 5 Data Change

Discussion of data changes for drift detection confirmation

## 6 Results and Demo

Review of all results and model monitoring drift detection demo

# Demo – Model Monitoring In Action!

The screenshot shows a Jupyter Notebook interface with the following details:

- Title Bar:** Notebooks > ml\_ops\_final\_project
- Section Header:** ml\_ops\_final\_project
- Description:** Add description
- Last Saved:** Last saved: 1 minute ago
- Code Cell:** 5b. Deploy the native XG-Boost model to AWS for inference
- Code Content:**

```
▶ ▾ ✓ 1 minute ago | Elapsed time: 3s
1 # Set S3 configurations for model saving
2 MODEL_KEY = 'model/model.tar.gz'
3 MODEL_S3_PATH = f's3://{BUCKET_NAME}/{MODEL_KEY}'
4 ENDPOINT_NAME = 'pokemon-model'
5
6 # Save the model
7 xgb_model.save_model('xgboost-model')
8
9 # Convert to tar.gz format (SageMaker requirement)
10 with tarfile.open('model.tar.gz', 'w:gz') as tar:
11     tar.add('xgboost-model')
12     print("Model packaged as model.tar.gz")
13
14 # Upload to S3
15 S3_CLIENT.upload_file(
16     Filename='model.tar.gz',
17     Bucket=BUCKET_NAME,
18     Key=MODEL_KEY
19 )
20 print(f"Model uploaded to: {S3_PATH}")
21
22 xgb_model_sm = XGBoostModel(
23     model_data= MODEL_S3_PATH,
24     role= ROLE_ARN,
25     framework_version= '1.7-1',
26     sagemaker_session= SESSION
27 )
28
29 print("\nDeploying model to SageMaker endpoint...")
30
```
- Kernel:** Python
- Resource Usage:** Python 3.11, 2 vCPU + 4 GiB, 4.00%, 53.99%, 2.00%
- Status:** Ready

# Drift Report Comparison: X\_test

1 Tests	1 Success	0 Warning	0 Fail	0 Error
All tests				
<b>Number of Drifted Features</b>				
 The drift is detected for 2 out of 21 features. The test threshold is lt=3.				Details
Feature name	Stattest	Drift score	Threshold	Data Drift
Type_2	K-S p_value	0.017	0.05	Detected
Has_Type_2	Z-test p_value	0.002	0.05	Detected
prediction	Z-test p_value	0.163	0.05	Not detected
Type_1	K-S p_value	0.562	0.05	Not detected
Total	K-S p_value	0.053	0.05	Not detected
Speed	K-S p_value	0.217	0.05	Not detected
Sp_Def	K-S p_value	0.08	0.05	Not detected
Sp_Atk	K-S p_value	0.468	0.05	Not detected
Height_m	K-S p_value	0.17	0.05	Not detected
Weight_kg	K-S p_value	0.545	0.05	Not detected
Has_Egg_Group_2	Z-test p_value	0.077	0.05	Not detected
Generation	K-S p_value	0.945	0.05	Not detected
Egg_Group_2	K-S p_value	0.545	0.05	Not detected
Egg_Group_1	K-S p_value	0.999	0.05	Not detected
Defense	K-S p_value	0.056	0.05	Not detected

# Drift Report Comparison: X\_test\_modified

Tests	Success	Warning	Fail	Error
All tests				
Number of Drifted Features				
prediction	Z-test p_value	0.0	0.05	Detected
Attack	K-S p_value	0.0	0.05	Detected
Type_2	K-S p_value	0.0	0.05	Detected
Type_1	K-S p_value	0.0	0.05	Detected
Total	K-S p_value	0.0	0.05	Detected
Sp_Def	K-S p_value	0.005	0.05	Detected
Has_Type_2	Z-test p_value	0.002	0.05	Detected
Speed	K-S p_value	0.217	0.05	Not detected
Sp_Atk	K-S p_value	0.468	0.05	Not detected
Height_m	K-S p_value	0.17	0.05	Not detected
Has_Egg_Group_2	Z-test p_value	0.077	0.05	Not detected
HP	K-S p_value	0.295	0.05	Not detected
Generation	K-S p_value	0.945	0.05	Not detected
Egg_Group_2	K-S p_value	0.545	0.05	Not detected
Enc_Genus_1	K-S p_value	0.000	0.05	Not detected

# GitHub Link:

<https://github.com/bhstoller/ml-ops-fp/tree/main>

# Questions?